

Introduction

- Moore's Law enabled:
 - Deep memory hierarchy
 - Wide SIMD units
 - Deep pipelines
 - Branch prediction
 - Out-of-order execution
 - Speculative prefetching
 - Multithreading
 - Multiprocessing



Whats the big deal ?

Figure 1. The historical virtuous cycle of universal processors (a) is turning into a fragmentation cycle (b).

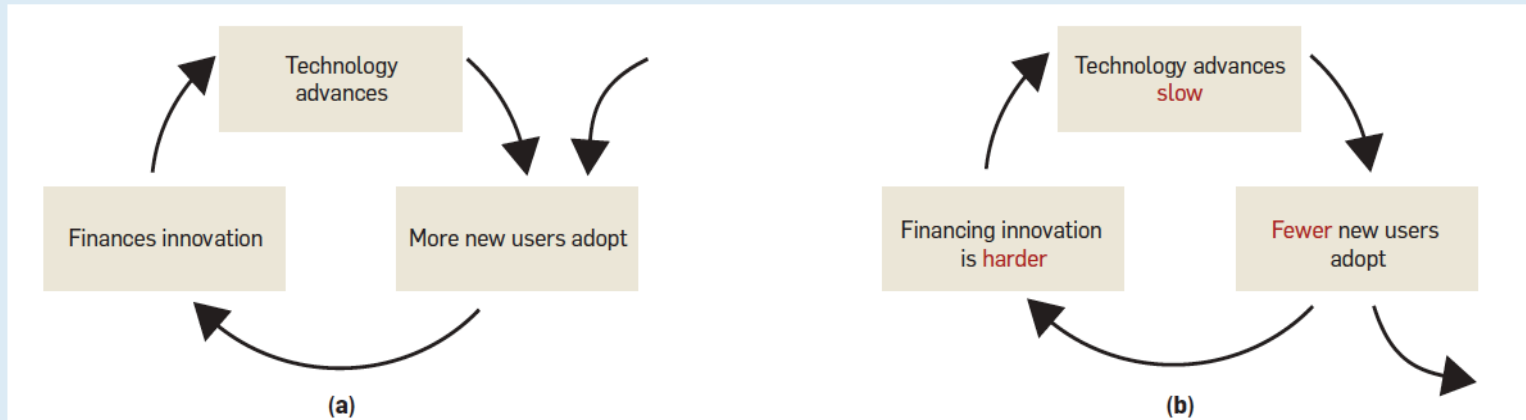




Figure 2. Rate of improvement in microprocessors, as measured by (a) Annual performance improvement on the SPECint benchmark, ^{7-appx} and (b) Annual quality-adjusted price decline. ^{1-appx}



Whats the big deal ?



DSA Philosophy 101.....

- Use dedicated memories to minimize data movement
- Invest resources into more arithmetic units or bigger memories
- Use the easiest form of parallelism that matches the domain
- Reduce data size and type to the simplest needed for the domain
- Use a domain-specific programming language



Examples from Textbook

Guideline	TPU	Catapult	Crest	Pixel Visual Core
Design target	Data center ASIC	Data center FPGA	Data center ASIC	PMD ASIC/SOC IP
1. Dedicated memories	24 MiB Unified Buffer, 4 MiB Accumulators	Varies	N.A.	Per core: 128 KiB line buffer, 64 KiB P.E. memory
2. Larger arithmetic unit	65,536 Multiply-accumulators	Varies	N.A.	Per core: 256 Multiply-accumulators (512 ALUs)
3. Easy parallelism	Single-threaded, SIMD, in-order	SIMD, MISD	N.A.	MPMD, SIMD, VLIW
4. Smaller data size	8-Bit, 16-bit integer	8-Bit, 16-bit integer 32-bit Fl. Pt.	21-bit Fl. Pt.	8-bit, 16-bit, 32-bit integer
5. Domain-specific lang.	TensorFlow	Verilog	TensorFlow	Halide/TensorFlow



Examples from Textbook

Guideline	TPU	Catapult	Crest	Pixel Visual Core
Design target	Data center ASIC	Data center FPGA	Data center ASIC	PMD ASIC/SOC IP
1. Dedicated memories	24 MiB Unified Buffer, 4 MiB Accumulators	Varies	N.A.	Per core: 128 KiB line buffer, 64 KiB P.E. memory
2. Larger arithmetic unit	65,536 Multiply-accumulators	Varies	N.A.	Per core: 256 Multiply-accumulators (512 ALUs)
3. Easy parallelism	Single-threaded, SIMD, in-order	SIMD, MISD	N.A.	MPMD, SIMD, VLIW
4. Smaller data size	8-Bit, 16-bit integer	8-Bit, 16-bit integer 32-bit Fl. Pt.	21-bit Fl. Pt.	8-bit, 16-bit, 32-bit integer
5. Domain-specific lang.	TensorFlow	Verilog	TensorFlow	Halide/TensorFlow




DNN's: Today's Driving Application

- Inspired by neuron of the brain
- Computes non-linear "activation" function of the weighted sum of input values
- Neurons arranged in layers

Name	DNN layers	Weights	Operations/Weight
MLP0	5	20M	200
MLP1	4	5M	168
LSTM0	58	52M	64
LSTM1	56	34M	96
CNN0	16	8M	2888
CNN1	89	100M	1750



Multi-Layer Perceptrons




$$\max(0, [O_{N-1,0} \ O_{N-1,1} \ O_{N-1,2}] \cdot \begin{bmatrix} W_{0,0} & W_{0,1} & W_{0,2} \\ W_{1,0} & W_{1,1} & W_{1,2} \\ W_{2,0} & W_{2,1} & W_{2,2} \end{bmatrix}) = [O_{N,0} \ O_{N,1} \ O_{N,2}]$$



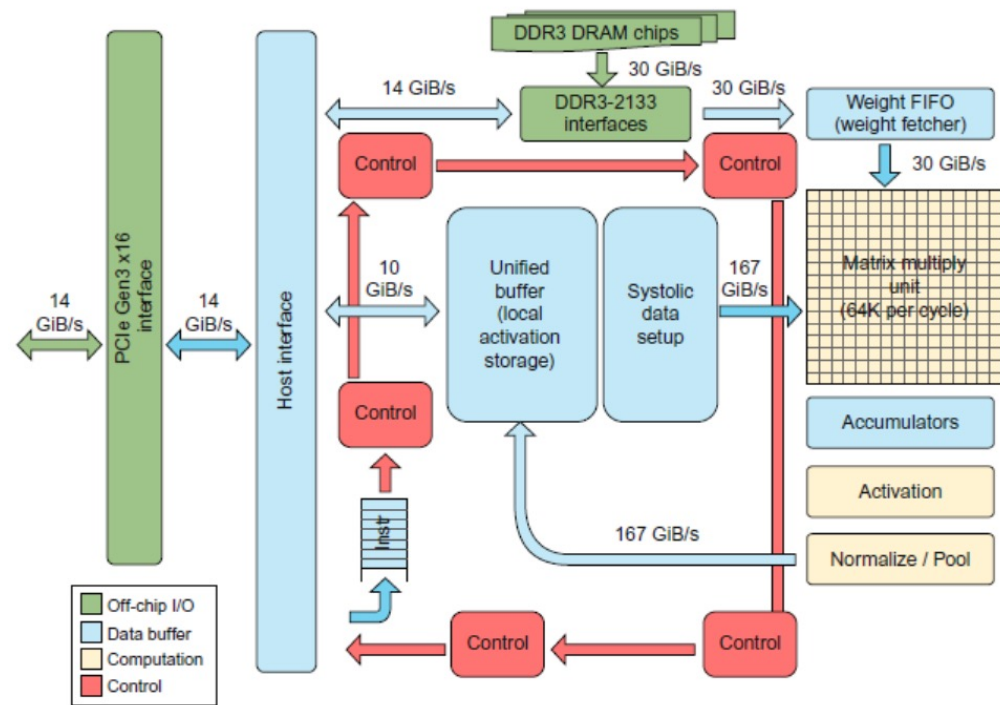
Multi-Layer Perceptrons

- Parameters:
 - $\text{Dim}[i]$: number of neurons
 - $\text{Dim}[i-1]$: dimension of input vector
 - Number of weights: $\text{Dim}[i-1] \times \text{Dim}[i]$
 - Operations: $2 \times \text{Dim}[i-1] \times \text{Dim}[i]$
 - Operations/weight: 2



Tensor Processing Unit

- Google's DNN ASIC
- 256×256 8-bit matrix multiply unit
- Large software-managed scratchpad
- Coprocessor on the PCIe bus


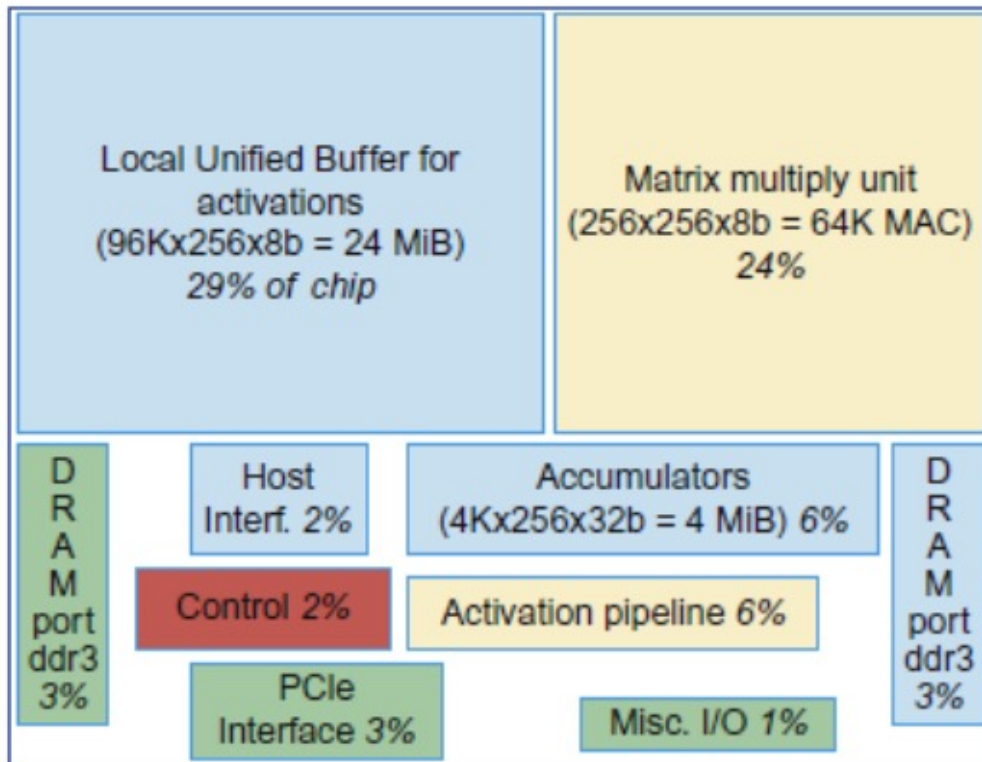


TPU ISA

- Read_Host_Memory
 - Reads memory from the CPU memory into the unified buffer
- Read_Weights
 - Reads weights from the Weight Memory into the Weight FIFO as input to the Matrix Unit
- MatrixMatrixMultiply/Convolve
 - Perform a matrix-matrix multiply, a vector-matrix multiply, an element-wise matrix multiply, an element-wise vector multiply, or a convolution from the Unified Buffer into the accumulators
 - takes a variable-sized $B \times 256$ input, multiplies it by a 256×256 constant input, and produces a $B \times 256$ output, taking B pipelined cycles to complete
- Activate
 - Computes activation function
- Write_Host_Memory
 - Writes data from unified buffer into host memory



TPU ISA

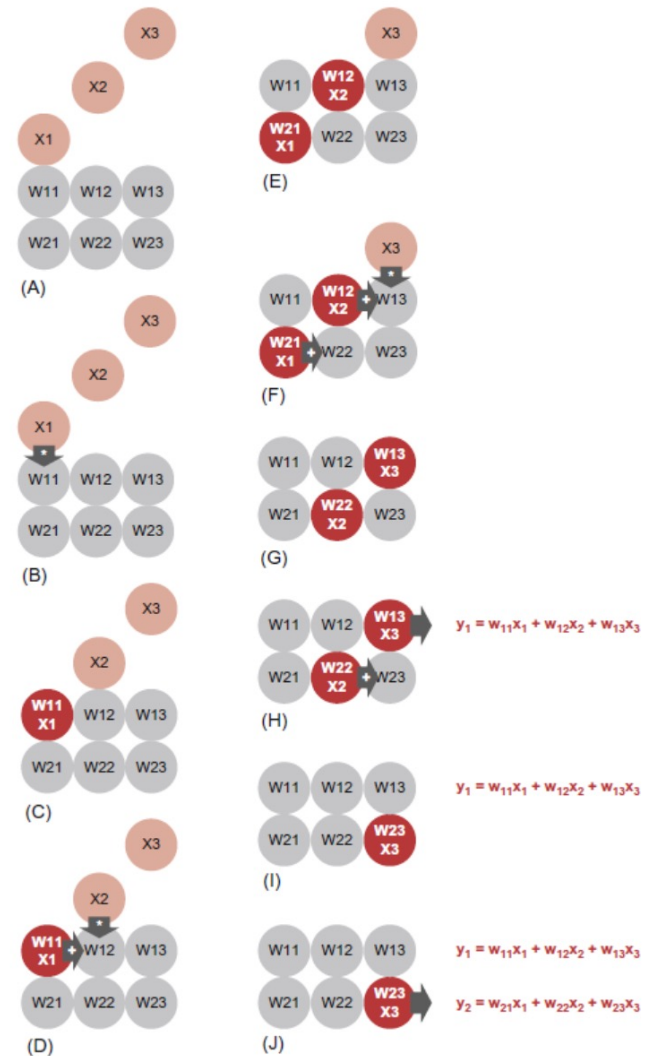



TPU ISA

- Matrix Multiply Unit (MPU) is a heart of TPU.
 - It contains 256x256 MACs(Multiply ACcumulate unit).
 - Performing 8-bit multiply and adds on signed or unsigned integers.
 - Output is 16-bit data.
 - 16-bit products are collected in the 4MiB of 32-bit accumulators.
 - The 4MiB represents 4096 node.
 - Each node has 256-element of 32-bit accumulators.
 - The matrix unit produces one 256-element partial sum per clock cycle
 - The matrix unit holds one 64KiB tile of weights plus one for double buffering. (To hide the 256 cycles it takes to shift a tile in)
 - Single weight is 8-Bit.



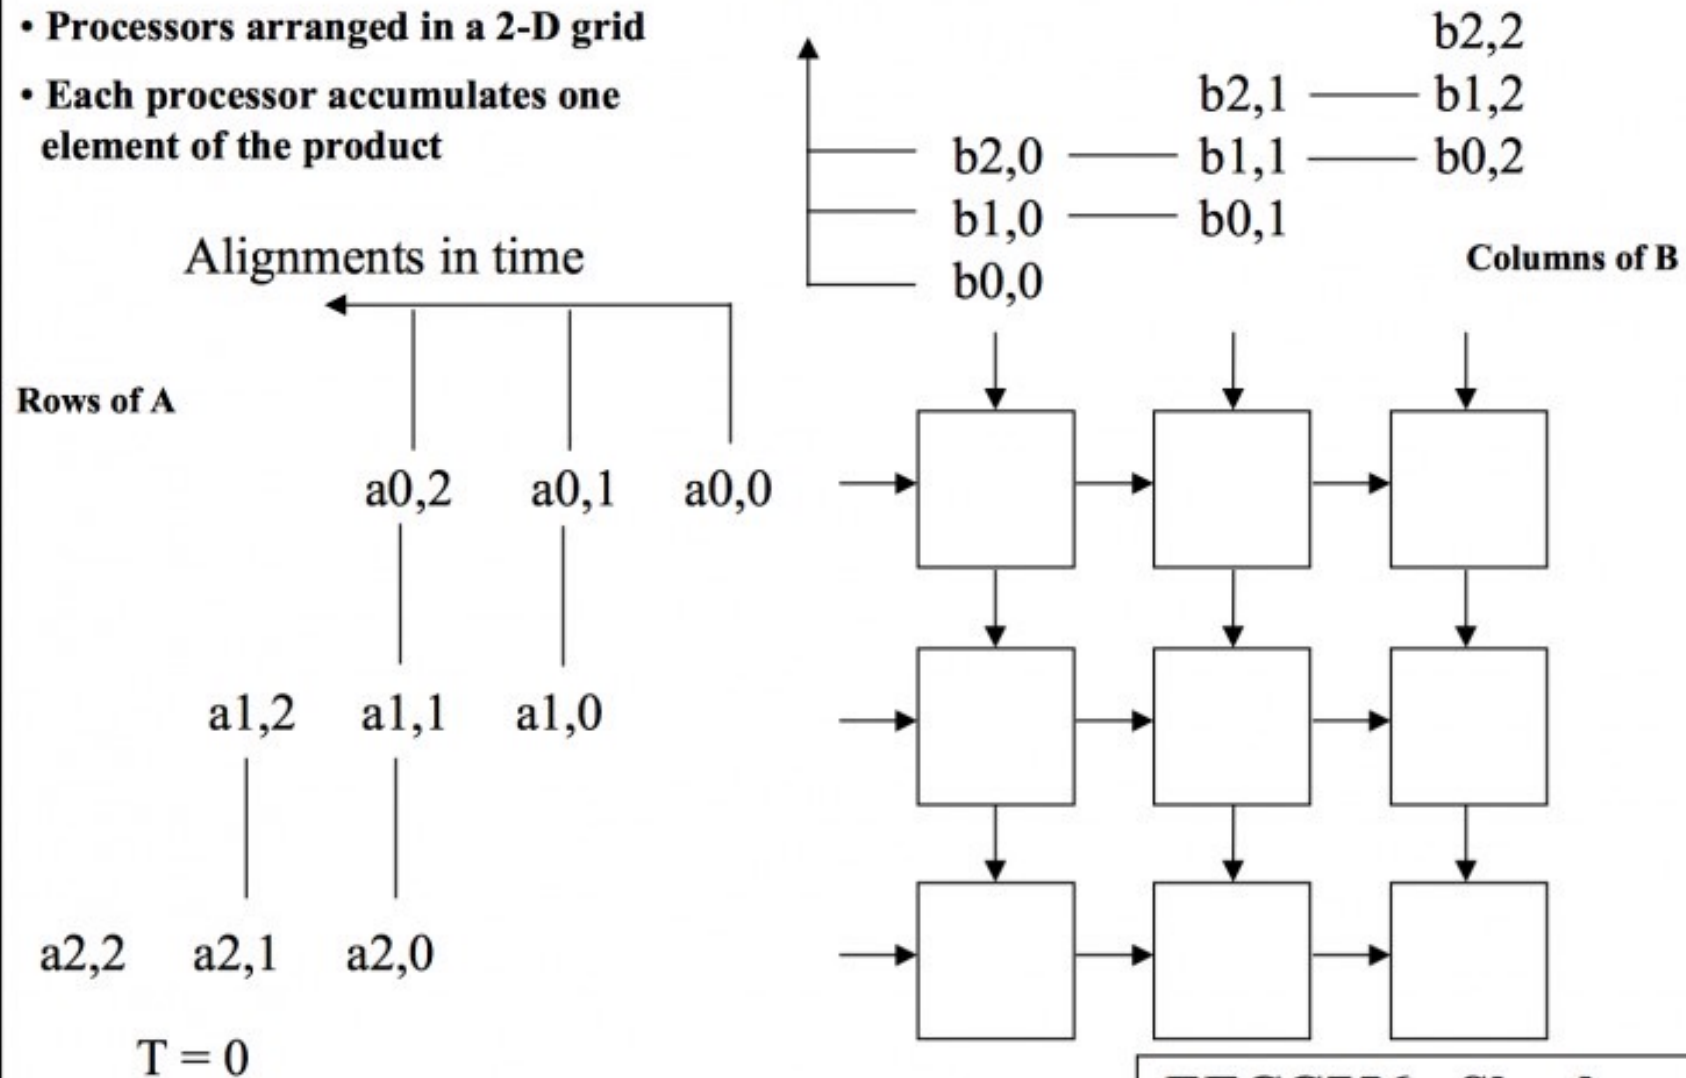
Structured as a Systolic Array



Systolic Array Example:

3x3 Systolic Array Matrix Multiplication


- Processors arranged in a 2-D grid
- Each processor accumulates one element of the product



EECC756 - Shaaban

Example source: <http://www.cs.hmc.edu/courses/2001/spring/cs156/>

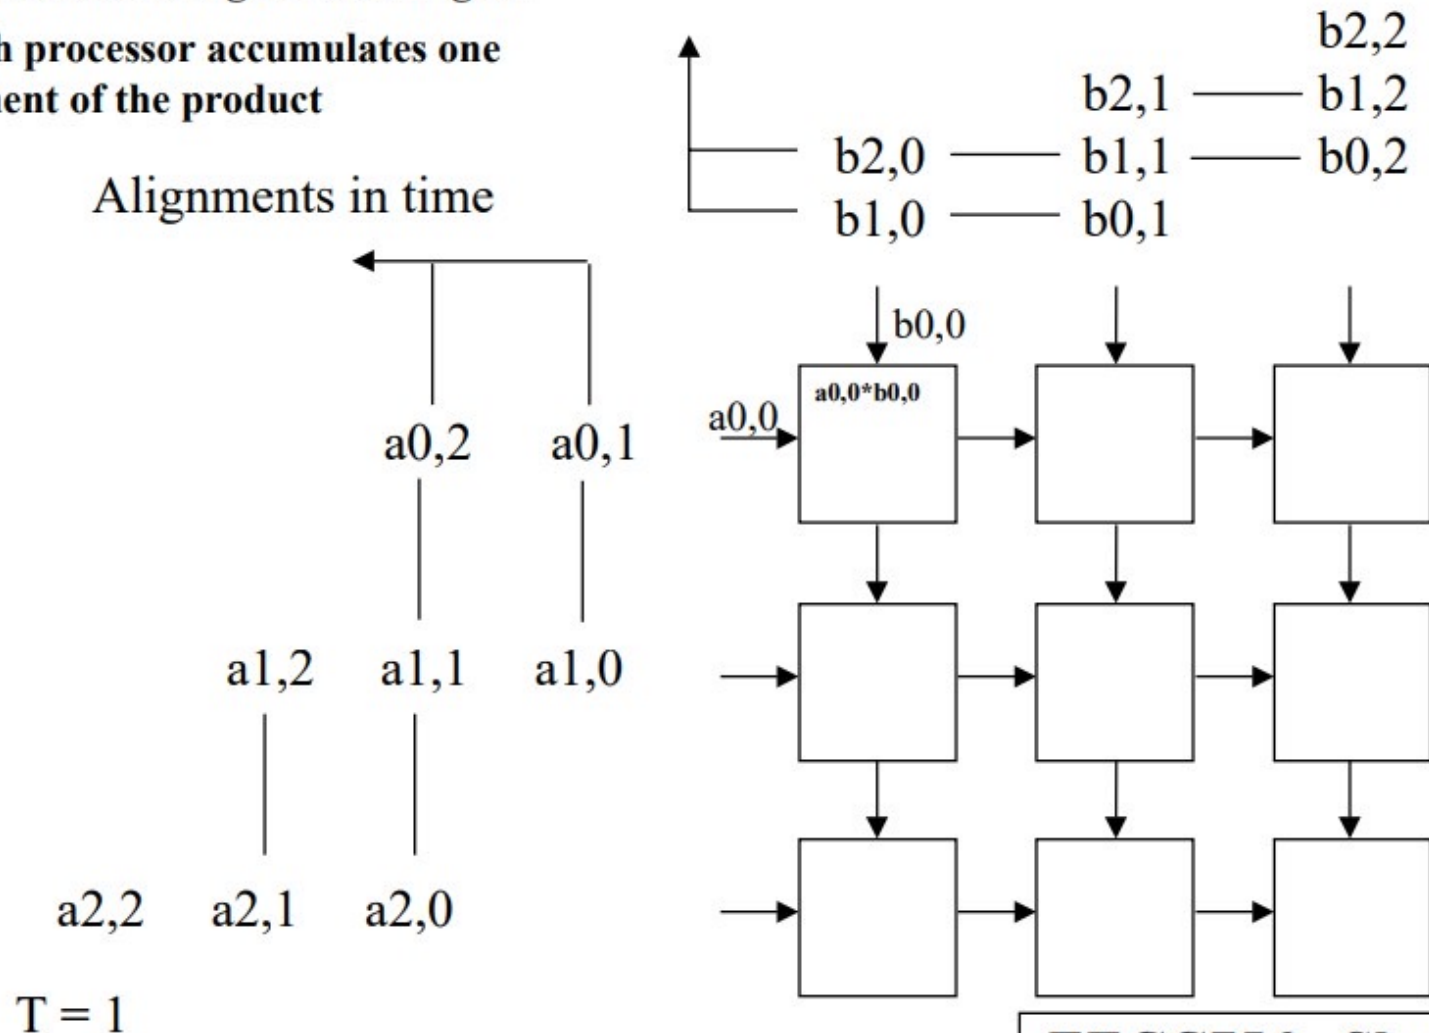
#2 lec # 1 Spring 2003 3-11-2003



Systolic Array Example:

3x3 Systolic Array Matrix Multiplication


- Processors arranged in a 2-D grid
- Each processor accumulates one element of the product



EECC756 - Shaaban

Example source: <http://www.cs.hmc.edu/courses/2001/spring/cs156/>

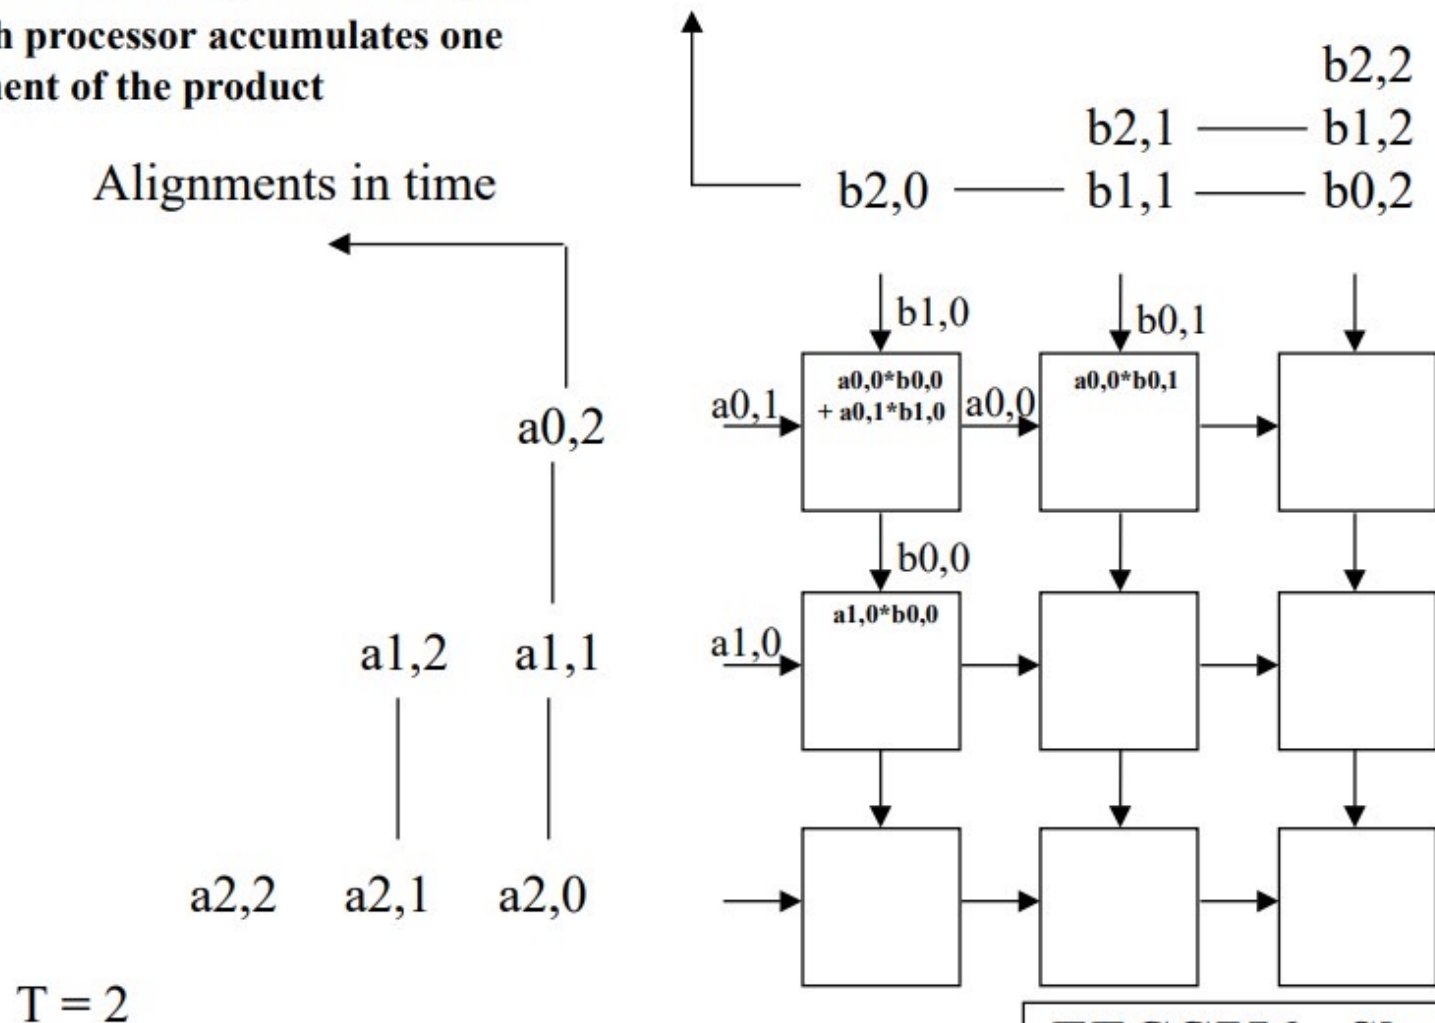
#3 lec #1 Spring 2003 3-11-2003



Systolic Array Example:

3x3 Systolic Array Matrix Multiplication


- Processors arranged in a 2-D grid
- Each processor accumulates one element of the product



EECC756 - Shaaban

Example source: <http://www.cs.hmc.edu/courses/2001/spring/cs156/>

#4 lec #1 Spring 2003 3-11-2003



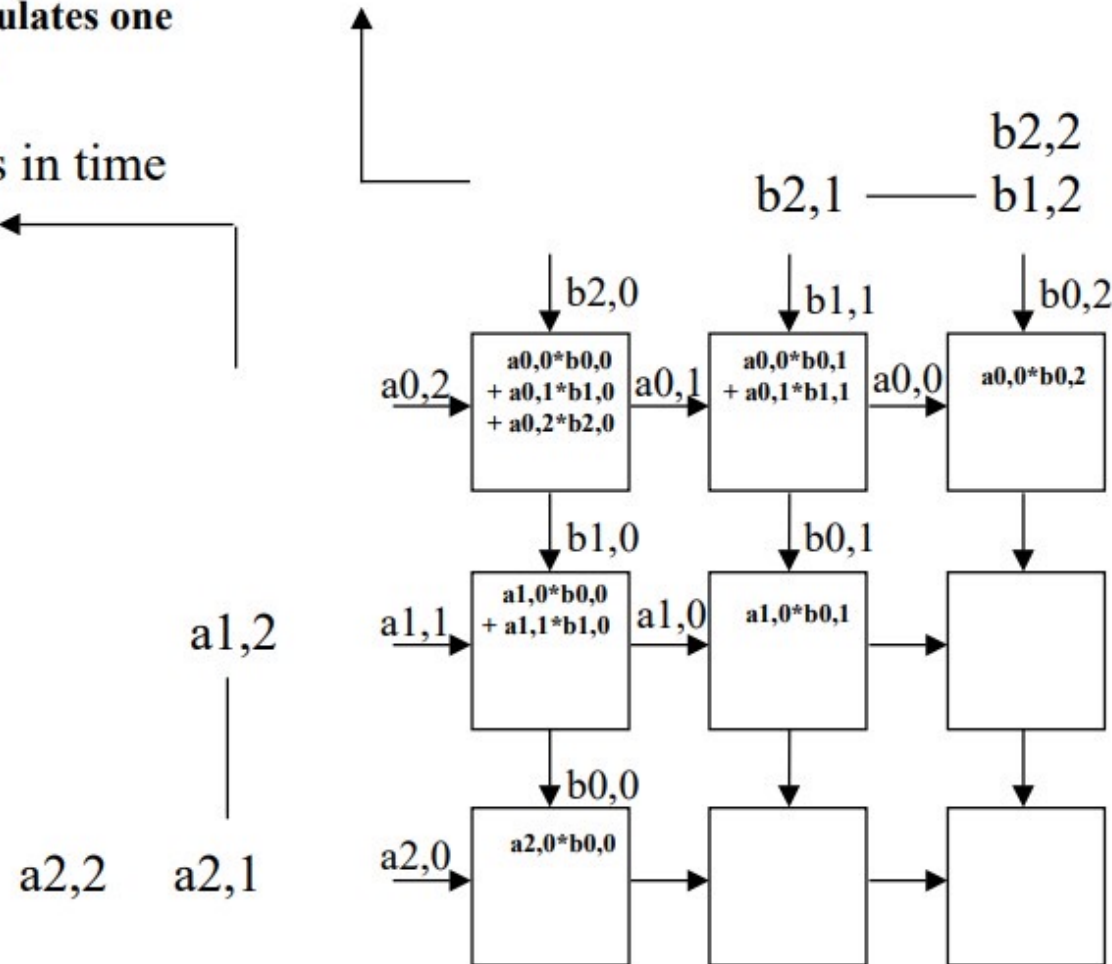
Systolic Array Example:

3x3 Systolic Array Matrix Multiplication

- Processors arranged in a 2-D grid
- Each processor accumulates one element of the product

Alignments in time

$T = 3$



EECC756 - Shaaban

Example source: <http://www.cs.hmc.edu/courses/2001/spring/cs156/>

#5 lec #1 Spring 2003 3-11-2003



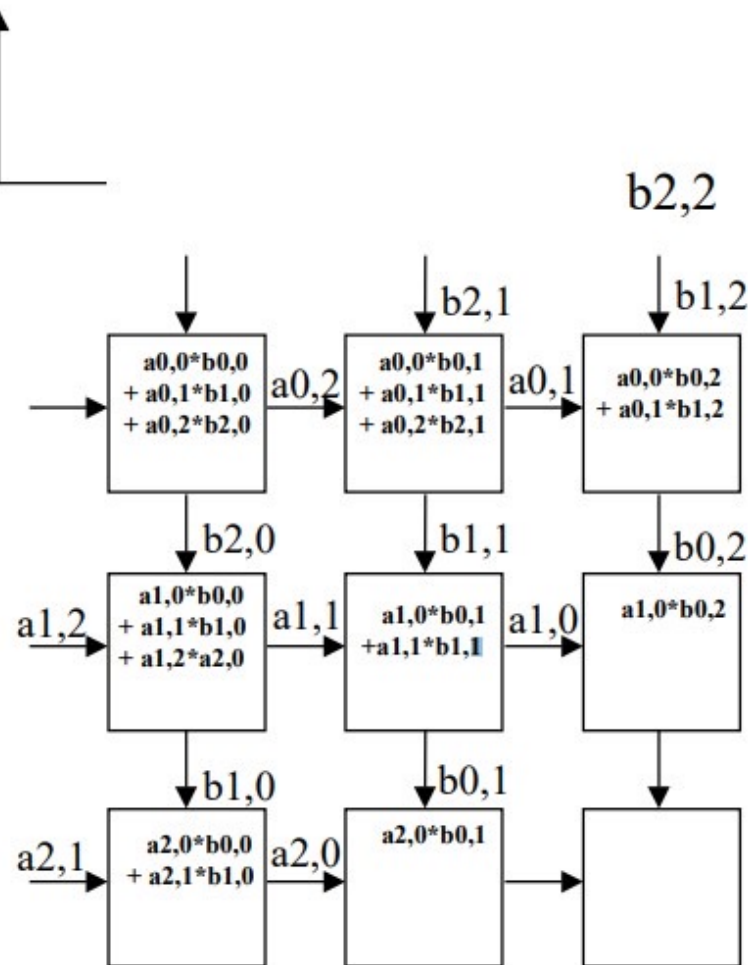
Systolic Array Example:

3x3 Systolic Array Matrix Multiplication

- Processors arranged in a 2-D grid
- Each processor accumulates one element of the product

Alignments in time


$T = 4$



EECC756 - Shaaban

Example source: <http://www.cs.hmc.edu/courses/2001/spring/cs156/>

#6 lec # 1 Spring 2003 3-11-2003

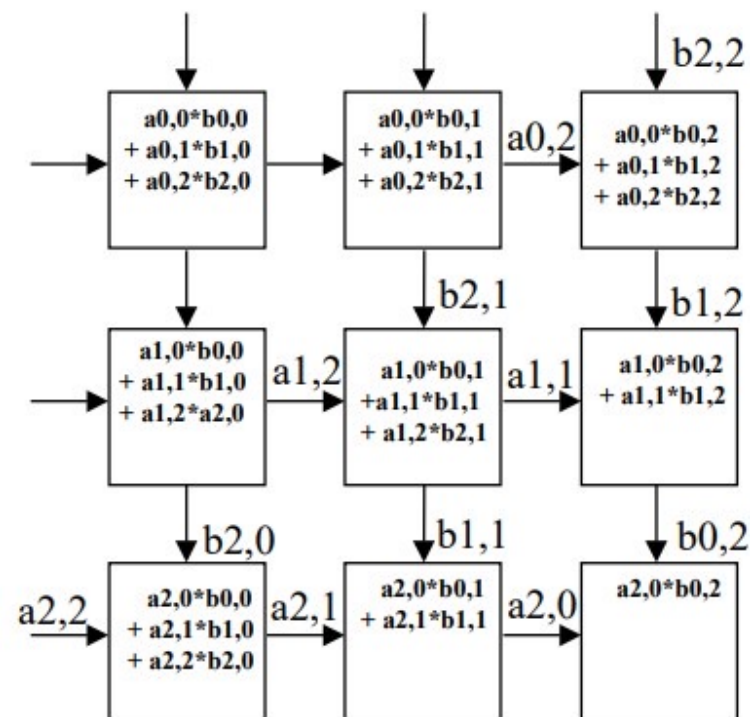


Systolic Array Example:

3x3 Systolic Array Matrix Multiplication

- Processors arranged in a 2-D grid
- Each processor accumulates one element of the product

Alignments in time




T = 5

EECC756 - Shaaban

Example source: <http://www.cs.hmc.edu/courses/2001/spring/cs156/>

#7 lec # 1 Spring 2003 3-11-2003

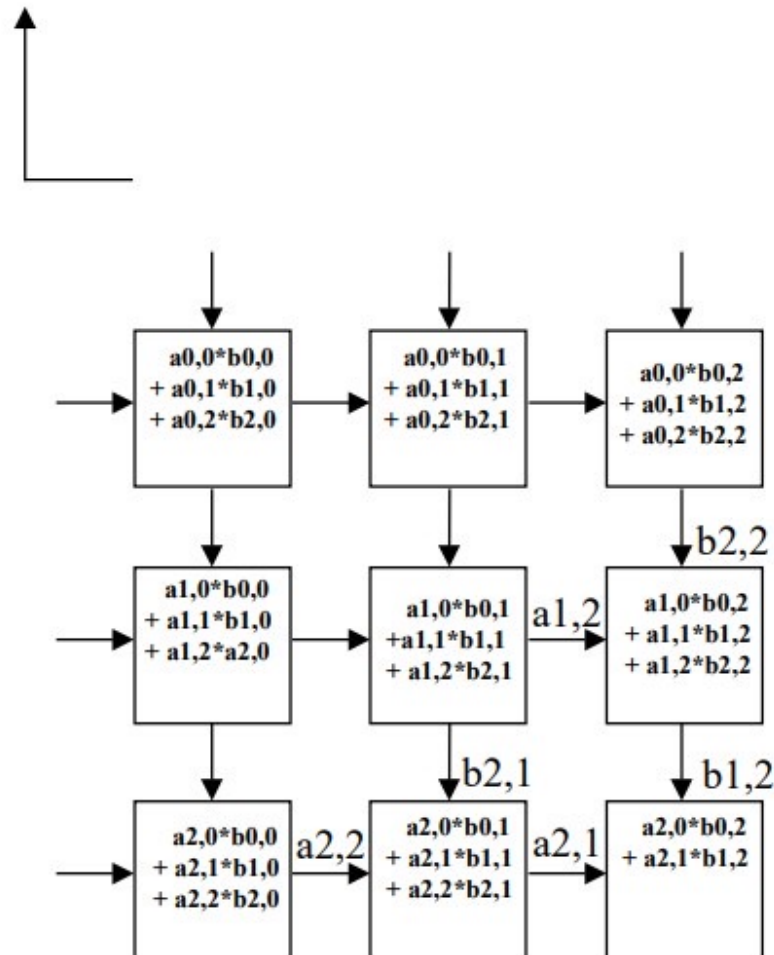


Systolic Array Example:

3x3 Systolic Array Matrix Multiplication

- Processors arranged in a 2-D grid
- Each processor accumulates one element of the product

Alignments in time




$T = 6$

EECC756 - Shaaban

Example source: <http://www.cs.hmc.edu/courses/2001/spring/cs156/>

#8 lec # 1 Spring 2003 3-11-2003



Systolic Array Example:

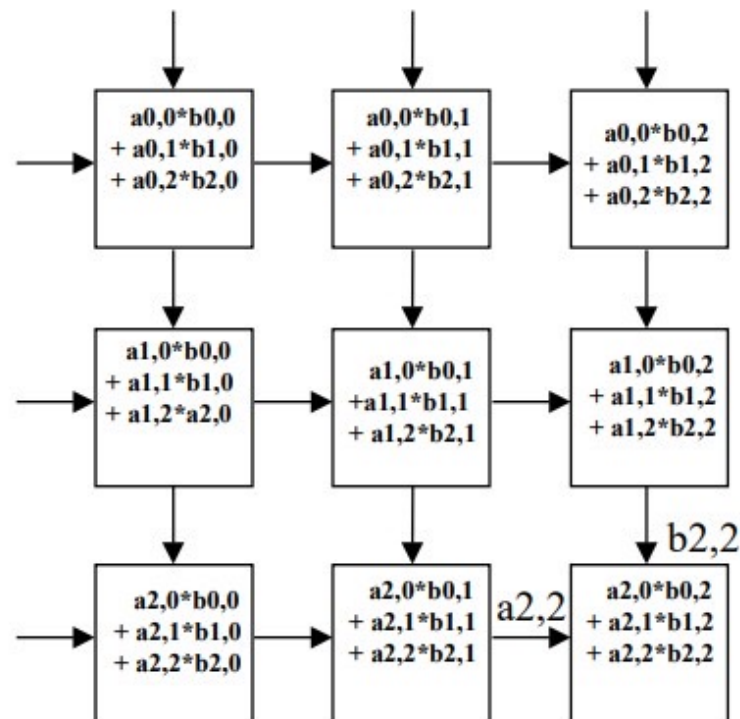
3x3 Systolic Array Matrix Multiplication

- Processors arranged in a 2-D grid
- Each processor accumulates one element of the product

Alignments in time

Done

$T = 7$



EECC756 - Shaaban

Example source: <http://www.cs.hmc.edu/courses/2001/spring/cs156/>

#9 lec #1 Spring 2003 3-11-2003



The TPU and the Guidelines

- Use dedicated memories
 - 24 MiB dedicated buffer, 4 MiB accumulator buffers
- Invest resources in arithmetic units and dedicated memories
 - 60% of the memory and 250X the arithmetic units of a server-class CPU
- Use the easiest form of parallelism that matches the domain
 - Exploits 2D SIMD parallelism
- Reduce the data size and type needed for the domain
 - Primarily uses 8-bit integers
- Use a domain-specific programming language
 - Uses TensorFlow

